

Assignment 03

Part II Programming (110 points)

Due: Beginning of the class, Mar 23rd

Specification: Multi-level Paging with TLB

You are given considerable latitude as to how you choose to design your data structures and interfaces. However, your implementation **must** fulfill the **implementation requirements** specified in the **Level** structure and **mandatory interfaces** below. A sample data structure is summarized below and advice on how it might be used is given on the Canvas assignment page. Refer to [pageTableDS.pdf](#).

Page table structures

- **PageTable** – Top level descriptor describing attributes of the N level page table and containing a pointer to the level 0 (root node) page tree/table structure.
 - PageTable captures your multi-level paging information (as put in the pagetableDS.pdf) which is used for every page Level (or node) object: number of levels, the bit mask and shift information for extracting the part of VPN pertaining to each level, the number of entries in each Level array to the next level objects, or the array to the PFN mapping for the leaf level / node, etc.
 - Since the tree operations start from root node, it would be convenient to have the PageTable have a reference / pointer to the root node (Level) of the page tree.
- **Level** (or PageNode) – An entry for an arbitrary level, this is the structure (or class) which represents one of the sublevels in the page tree/table.
 - Level is essentially the structure for the multi-level page tree node. The multi-level paging is about splitting and storing the VPN information into a tree data structure along the tree paths. Starting from the root node, each tree path (the root node to a leaf node) captures the mapping from a VPN to a PFN, and the mapped PFN would be stored at the leaf node level.
 - It contains an array of either pointers to the next level or page mappings.
 - For non-leaf or interior level nodes, an array of Level* (or PageNode*) pointers to the next level, it is essentially a double Level**
 - For leaf level nodes, an array of mappings, each mapping maps a VPN to a physical frame
 - Note you can have both an Level* array (for interior nodes) and a mapping array (for leaf nodes) included in the Level structure and use the one dependent on where the level is at: i.e., interior or leaf level.
 - **Implementation requirements (violation would incur 50% penalty of autograding):**
 - You **MUST** implement this Level / PageNode structure as a tree.
 - **You must NOT use a hash map for storing children levels / nodes**, you could use either Level** or std::vector<Level*> or similar **array** representation.
- **Map** – A structure containing information about the mapping of a page to a frame, used in leaf nodes of the tree.

Page table mandatory interfaces

Implementation requirements (violation would incur **50% penalty** of autograding): You **must** implement similar functions for multi-level paging as proposed below. Your exact function signatures may vary, but the functionality should be the same.

All other interfaces may be developed as you see fit.

- `unsigned int virtualAddressToVPN(unsigned int virtualAddress, unsigned int mask, unsigned int shift)`
 - Given a virtual address, apply the given bit mask and shift right by the given number of bits. Returns the virtual page number. This function can be used to access the **virtual page number of any level** by supplying the appropriate parameters.
 - Example: Suppose the level two pages occupied bits 22 through 27, and we wish to extract the second level page number of address 0x3c654321. `virtualAddressToVPN(0x3c654321, 0x0FC00000, 22)` should return 0x31 (decimal 49). Remember, this is computed by taking the bitwise **and** operation of 0x3c654321 and 0x0FC00000, which is 0x0C400000. We then shift right by 22 bits. The last five hexadecimal zeros take up 20 bits, and the bits higher than this are 1100 0110 (C6). We shift by two more bits to have the 22 bits, leaving us with 11 0001, or 0x31.
 - Check out the given bitmasking-demo.c for bit masking and shifting for extracting bits in the hexadecimal number.
 - **Note:** to get the full Virtual Page Number (VPN) from all page levels, you would construct the bit mask for all bits preceding the offset bits, take the bitwise **and** of the virtual address and the mask, then shift right for the number of offset bits.
VPN from all levels combined is needed for caching the Virtual Page Number (VPN) to Physical Frame Number (PFN) mapping in the TLB, see below.
- `Map * lookup_vpn2pfn(PageTable *pageTable, unsigned int virtualAddress)`
 - Given a page table and a virtual address, return the mapping of virtual address to physical frame (in Map*) from the page table. You must have an appropriate return value for when the virtual page is not found (e.g. NULL if this is the first time the virtual page has been seen). Note that if you use a different data structure than the one proposed, this may return a different type, but the function name and idea should be the same. Similarly, If `lookup_vpn2pfn` was a method of the C++ class `PageTable`, the function signature could change in an expected way: `Map * PageTable::lookup_vpn2pfn(unsigned int virtualAddress)`. This advice should be applied to other page table functions as appropriate.
- `void insert_vpn2pfn(PageTable *pagetable, unsigned int virtualAddress, unsigned int frame)`
 - Used to add new entries to the page table when we have discovered that a page has not yet been allocated (`lookup_vpn2pfn` returns NULL). Frame is the frame index which corresponds to the page number of the virtual address. Use a frame number of 0 the first time this is called, and increment the frame index by 1 each time a new page→frame map is needed. If you wish, you may replace void with int or bool and return an error code if unable to allocate memory for the page table. **HINT:** If you are inserting a page, you do not always add nodes at every level. The Map structure may already exist at some or all of the levels.

Translation Lookaside Buffer (TLB)

Important: You may use any C++ standard template library class for implementing the TLB.

As part of the MMU translation simulation, you are required to implement a simulation of TLB for caching the virtual page to physical frame mappings from the page table, each TLB entry caches one mapping of a Virtual Page Number to a Physical Frame Number (**VPN → PFN**, whichever the way you prefer to implement).

The TLB size (max number of mappings or entries) can be designated by an optional command line argument (see -c in the user interface specification below). The default size is **0** if the argument is not specified.

Your TLB cache will use an approximation of the least recently used (LRU) cache replacement policy. Least Recently Used (LRU) caching scheme evicts the least recently used cache entry when cache is full to make room for bringing in a new entry that is referenced and not in the cache. In doing so, it needs to track the most recent access to the cache entries to allow quick identification of which entry is the least recently used. Specifically, when TLB is full (number of mappings cached reaches the TLB size) and a TLB miss occurs, you would first go to the page table (or demand paging to allocate a new frame for mapping) to find the mapping, evict an existing mapping from the TLB using LRU replacement policy, and then you can insert the new mapping to the TLB.

Here you would use an approximation of the least recently used policy which tracks a set of recently accessed **pages** to evict the oldest accessed one for cache replacement. It is reasonable to track a small set of recently accessed pages for approximating the perfect LRU cache policy. In this assignment:

- You would track the **8** (or the number of current TLB entries if it is less than **8**) **most recently accessed pages (with distinct VPNs)** and their access times for implementing the LRU cache replacement.

Note:

- A perfect LRU cache policy would track the access times of all cached pages in TLB which would add a lot more overhead and usually is not practical.

Important:

- After a cache miss, whether it is a page hit or page miss, the currently accessed page always needs to be inserted in the TLB. And when TLB is full, you would need to perform cache replacement.
- As part of the implementation for cache replacement, whether it is a cache hit or miss, you always need to update the most recently accessed pages.
- With every virtual address access, insert or update the recently accessed pages for the virtual page (that contains the virtual address) being accessed and the access time. If the page being accessed is not in the recently accessed pages and if the number of recently accessed pages reaches **8**, the oldest page (the page with the smallest last access time) tracked there should be removed, and the page being accessed should then be added to the recently accessed pages.
- It is very **important** to understand that you need to track the **8 (or number of current cache entries) most recently accessed PAGES (with distinct VPNs)**, NOT the last **8** accessed addresses. There could often be a series of access attempts to the virtual addresses of the **same** page.
- When evicting a TLB entry due to cache replacement, make sure the entry with the same VPN (oldest page tracked) in most recently accessed pages is also removed.
- You may want to use a virtual time instead of the actual time for tracking each memory / page access time. An **address count** would be a good choice serving as a virtual time, it starts from zero before the access of the first address and is incremented by 1 after reading / accessing each address from the trace file.